

I will break down the fundamental, most critical concepts that underpin glibc heap security and exploitation into a simpler framework.

The complexity essentially boils down to three core areas: the structure of the memory blocks, the system for recycling freed blocks, and how these mechanisms are hijacked to gain control.

---

# 1. The Core Vulnerability: Data Overlap

The foundational issue that enables almost all heap exploitation lies in the dual nature of the memory blocks, known as **chunks**, which are allocated by malloc.

## The Dual-Purpose Chunk

When your program calls malloc(N), the heap manager finds a chunk large enough to hold N **bytes** of user data plus a small **metadata header** (containing information like the chunk's total size and status flags).

State	Memory Use	Risk Profile
Allocated	The memory region (starting after the header) is entirely available for user data.	The programmer is responsible for not writing past the end of this region (Buffer Overflow).
Freed	The heap manager <b>repurposes</b> the user data space to store its internal linked list pointers: the Forward Pointer (fd) and the Backward Pointer (bk).	If an attacker retains a pointer to this now-freed memory ( <b>Use-After-Free</b> ) or can overflow an adjacent block, they are now directly overwriting the heap manager's plumbing (fd and bk).

## The Exploit Hook

This **data overlap** means that controlling the contents of a *freed* chunk allows an attacker to control the pointers the allocator uses to manage its lists. By manipulating these pointers, an attacker can trick the allocator into returning a pointer to an arbitrary memory location, such as a function pointer or configuration data.

---

## 2. The Recycling System: The Four Bins

The glibc allocator (ptmalloc2) categorizes and manages freed chunks using a complex system of linked lists called **Bins**. This system is optimized for speed, and its rules create the primary targets for exploitation.

Bin Type	Purpose (The Strategy)	Structure & Order	Key Exploitable Trait
Fast Bins	<b>Maximal Speed:</b> Holds very small chunks (up to 88 bytes) intended for immediate reuse. They are never merged with adjacent free chunks.	<b>Singly Linked List (LIFO).</b> The last chunk freed is the first chunk reused. <sup>1</sup>	The simple singly-linked structure makes it susceptible to <b>double-free</b> attacks and <b>chunk forging</b> (like Fastbin Dup). <sup>3</sup>
Small Bins	<b>Balanced Speed:</b> Holds small chunks of fixed sizes (up to ~504 bytes). Unlike Fastbins, these chunks <b>are</b> merged when freed to control fragmentation. <sup>4</sup>	<b>Doubly Linked List (FIFO).</b> Insertion at the head, removal from the tail. <sup>4</sup>	The integrity checks required for managing doubly-linked lists (fd and bk) are targeted by attacks like <b>House of Lore</b> . <sup>5</sup>
Unsorted Bin	<b>Optimization Cache:</b> A single, temporary holding area where newly freed, consolidated Small and Large chunks are placed before being sorted into their correct bins. <sup>6</sup>	<b>Single Doubly Linked List.</b> The allocator checks this bin first, providing a "second chance" for quick reuse. <sup>6</sup>	Corruption allows for the <b>Unsorted Bin Attack</b> , which can leak addresses or write data by exploiting the delayed sorting process. <sup>3</sup>
Large Bins	<b>Large Storage:</b> Holds chunks larger	<b>Doubly Linked List, sorted by size.</b>	Targeting these lists provides access to

	than Small Bins, categorized by size ranges (not fixed sizes).	Allocation is slower as it requires traversing the list to find the best fit. <sup>6</sup>	large memory regions for overwriting and memory manipulation. <sup>3</sup>
--	--	--	--

### 3. The Goal: Arbitrary Control Primitives

All heap exploitation techniques aim to convert an underlying memory corruption bug (like a buffer overflow or use-after-free) into a powerful **primitive**: a reliable ability to read or write memory at any location.

#### Primitive 1: Arbitrary Allocation

- **Goal:** Force malloc() to return a pointer to an attacker-chosen address (P).
- **How:** By corrupting the fd pointer of a freed chunk to point to an address near P. The next time malloc is called for that size, it follows the corrupted link and returns a chunk that starts at P.
- **Key Example: House of Force** (Corrupting the **Top Chunk**). The top chunk is the largest block available. An attacker overwrites its size field with a gigantic value (e.g., -1). This tricks the allocator into believing the chunk spans nearly the entire virtual address space. By making a mathematically calculated request size, the attacker forces the top chunk pointer to "jump" precisely to a target address (P), which is then returned by a final malloc call.<sup>5</sup>

#### Primitive 2: Arbitrary Write

- **Goal:** Write a specific value to an attacker-chosen address.
- **How:** By exploiting the routines that manage the doubly-linked lists.
- **Key Example: Unlink Exploit** (Targeting Small/Unsorted Bins). When a chunk is removed (unlinked) from a doubly-linked list, the allocator performs critical pointer housekeeping: it updates the fd and bk pointers of the neighboring chunks to bypass the victim.<sup>7</sup> By corrupting the victim chunk's fd and bk pointers *before* it is unlinked, the attacker can redirect these housekeeping updates, resulting in arbitrary writes to memory.<sup>3</sup>

In essence, heap exploitation is the process of identifying a bug that allows memory corruption, and then meticulously engineering the heap's state to ensure that the allocator's internal logic operates on the corrupted metadata, ultimately delivering an arbitrary memory primitive.