

# **The Unitized Threat: Comprehensive Analysis and Countermeasures for Systemd Persistence (MITRE ATT&CK T1501)**

## **Section 1: Contextualizing Systemd as an Attack Surface**

Systemd represents the modern architecture for Linux system and service management, having replaced older initialization systems such as SysVinit and Upstart across many contemporary distributions.<sup>1</sup> This powerful suite of daemons and utilities centrally controls vital system resources, background processes, and overall system state.<sup>2</sup> Due to its inherent authority over the system lifecycle and ability to manage processes that persist across reboots, systemd is a critical and high-value target for adversaries seeking to maintain access to compromised endpoints.<sup>3</sup>

### **1.1 The Ubiquity and Authority of Systemd**

As the foundational init system, systemd utilizes standardized building blocks called units to define how resources behave.<sup>2</sup> These units encapsulate services, network resources, devices, mounts, and resource pools, enabling granular control over nearly every component of a Linux server. The centralized management provided by systemd allows threat actors, once they obtain the necessary privileges, to establish persistence that is resilient to system reboots or even credential changes.<sup>3</sup>

### **1.2 Mapping Systemd Persistence to MITRE ATT&CK**

The method of leveraging systemd for persistent access is cataloged within the MITRE ATT&CK framework as T1501: Systemd Service, primarily falling under the Persistence tactic (TA0003).<sup>5</sup> Specifically, the creation or modification of unit configuration files to execute malicious payloads is tracked under the sub-technique T1543.002: Create or Modify System Process: Systemd Service.<sup>1</sup>

However, the analysis of persistence techniques often necessitates considering synergistic execution methods. While T1543.002 addresses the service definition itself, systemd offers diverse activation mechanisms that intersect with other techniques. For instance, persistence

established via systemd timers often links directly to T1053 (Scheduled Task/Job).<sup>3</sup> Furthermore, if an adversary attempts to inject malicious code into an already running service process, this action may overlap with T1055 (Process Injection).<sup>7</sup>

### 1.3 Anatomy of a Systemd Unit File and Critical Directories

Systemd management relies entirely on unit configuration files, identifiable by file extensions such as `.service`, `.timer`, `.path`, and `.socket`.<sup>2</sup> These files dictate the behavior and life cycle of the resource they manage. Adversaries target these files for modification or creation, with specific locations depending on the desired scope of persistence:

- **System-level units:** These units typically require root access to be created or modified and are stored in global directories, most commonly `/etc/systemd/system/` and occasionally in system library directories like `/usr/lib/systemd/system/`.<sup>1</sup> Persistence established here grants the highest level of privilege, often running as the root user.
- **User-level units:** This is a crucial, often overlooked, vector. Any user can establish persistence without requiring root privileges by placing unit files in their user configuration directory: `$HOME/.config/systemd/user/`.<sup>1</sup>

A robust defense strategy must recognize that the tactical classification of T1501 frequently requires correlating the persistence mechanism (the `.service` unit, T1543.002) with its activation method (e.g., a `.timer` unit, T1053). If an adversary deploys a malicious service unit and pairs it with a timer unit to control its execution schedule, security analysts must monitor indicators for both the scheduling action and the payload delivery mechanism simultaneously. This integrated view ensures that indicators are not missed by treating the timer and the service as isolated events.

## Section 2: Primary T1501 Persistence Mechanisms: Service Unit Manipulation

The most direct and common method for achieving T1501 persistence involves manipulating the configuration directives within the ```` section of a standard `.service` unit file. These directives control when and how commands are executed throughout the service's operational lifecycle.

### 2.1 Leveraging Core Execution Directives

The foundational directive used for command execution upon service launch is `ExecStart`, which specifies the commands run when the service is started, either manually via `systemctl start` or automatically upon system boot.<sup>1</sup> Adversaries can achieve persistence by creating an entirely new, malicious service unit file, or by modifying the `ExecStart` directive of an existing, legitimate service to execute a payload.<sup>1</sup>

The creation and deployment of these malicious configurations are often automated by sophisticated offensive frameworks. For example, modules in toolkits like Metasploit can

handle the complex details of writing the malicious payload to disk, generating the corresponding unit file, and scheduling it for execution at the next system boot.<sup>4</sup> Once a unit file is written to disk, the command `systemctl enable <name>` ensures the service loads at boot, and `systemctl start <name>` executes the specified commands immediately.<sup>4</sup>

## 2.2 Advanced Execution Flow Manipulation for Redundancy

High-fidelity threat actors utilize auxiliary execution directives to maximize their opportunities for code execution and ensure resilience, even if the primary service lifecycle is interrupted. This technique aligns conceptually with the MITRE ATT&CK tactic of Redundant Access (T1108).<sup>7</sup>

Adversaries leverage execution hooks to stage payloads or launch secondary processes outside the main service process, potentially evading monitoring focused solely on the primary process ID (PID) defined by `ExecStart`.

- **Pre/Post Hooks:** The `ExecStartPre` directive executes commands *before* the main service starts, and `ExecStartPost` executes commands *after* it starts.<sup>1</sup> A sophisticated attacker might use `ExecStartPre` to launch a low-observable, memory-resident implant before the main service ever begins, ensuring access even if the main service is immediately terminated.
- **Stop and Clean-up Evasion:** The `ExecStop`, `ExecStopPre`, and `ExecStopPost` directives execute when a service is stopped.<sup>1</sup> This provides a potent vector for persistence resilience. An adversary can configure `ExecStopPost` to automatically restart the malicious payload or initiate a secondary persistence mechanism. If a security analyst attempts to stop the malicious service, the stop command itself triggers the relaunch, making traditional cleanup efforts ineffective.
- **Configuration Triggered Persistence:** The `ExecReload` directive executes when the service receives a configuration reload signal.<sup>4</sup> This allows an adversary to execute code unexpectedly any time an administrator attempts to refresh service configurations, providing an interactive, yet stealthy, trigger for persistence.

The exploitation of pre/post hooks signifies a robust, multi-stage persistence mechanism designed for survivability, not merely initial execution. If security monitoring focuses only on the primary service execution (`ExecStart`), the actor's actions using `ExecStartPre` or `ExecStopPost`—which launch secondary payloads—will be missed. This necessitates intercepting the configuration modification event itself, as the service configuration dictates the malicious behavior across all service states (start, stop, and reload).

## 2.3 Privilege Escalation Vectors via Service Configuration

Beyond simple persistence, `systemd` services offer configuration options that facilitate privilege escalation. The `User` directive specifies the user the service process should run as.<sup>1</sup> If an adversary achieves temporary root access (e.g., during an initial exploit phase), they can configure the service unit file to run perpetually as the root user upon every subsequent boot. This effectively institutionalizes root access, fulfilling the privilege escalation component

frequently associated with T1543.002.<sup>1</sup>

Additionally, threat actors may use symbolic links within the critical unit directories (e.g., /etc/systemd/system/) to point to the actual malicious payload located elsewhere on the filesystem.<sup>1</sup> This technique helps bypass simple directory monitoring by security tools, as the payload itself resides in an unexpected or less-audited location, while systemd remains configured to find the payload using the symbolic link.

## Section 3: Obfuscated Persistence via Advanced Systemd Unit Types

Systemd’s architecture is modular, supporting several unit types beyond the standard .service file. Adversaries leverage these advanced unit types—timers, paths, and sockets—to create event-driven persistence mechanisms that are often less understood and less actively audited than traditional services or cron jobs.

### 3.1 Timer Units: The Modern Scheduled Task (T1053 Intersection)

Systemd timers (.timer files) act as an integrated and highly advanced scheduling system, controlling the execution of associated service units.<sup>6</sup> While functionally similar to the traditional cron system, timers offer significant technical advantages for resilient execution.

The superiority of systemd timers stems from several factors. Unlike cron, where logging requires manual configuration, timers integrate detailed log and status reports via journalctl<sup>10</sup>, which aids attackers in verifying successful execution. Critically, timers support dependency management, allowing the adversary to define required units (e.g., ensuring network services are active) before the payload service runs, preventing noisy failures.<sup>10</sup> Furthermore, systemd time events offer higher accuracy (second precision) and more flexible scheduling using sophisticated calendar specifications or monotonic times (e.g., relative to system boot or unit activation) compared to crontab's rigid structure.<sup>11</sup>

Table 1: Comparative Analysis of Persistence Scheduling (Cron vs. Systemd Timers)
Feature
Auditability/Logging
Dependency Management
Scheduling Format

### 3.2 Event-Driven Persistence via Path and Socket Units

Path and socket units facilitate a tactical evolution from time-based scheduling to *event-based activation*, which can significantly reduce the window of detection by keeping the malicious service dormant until a specific system event occurs.

- **Path Units:** Path units (.path files) leverage the kernel's inotify mechanism to trigger a service based on filesystem events.<sup>2</sup> This permits reactive persistence. An attacker can configure a .path unit to launch a malicious service whenever a specific critical file (such as a system log or configuration file) is accessed, modified, or created. The malicious code remains inactive and unseen until the triggering I/O event occurs.<sup>2</sup>
- **Socket Units:** Socket units (.socket files) enable socket-based activation, delaying the start of the service until traffic hits a specific network port or Inter-Process Communication (IPC) mechanism.<sup>1</sup> A malicious payload can be registered as a socket unit. When a user or remote system connects to the designated port, the systemd service (containing the payload) is dynamically launched. This conserves system resources and delays execution until interaction occurs, hindering discovery during routine process checks.
- **D-Bus Service Hijacking:** In highly sensitive systems, D-Bus service unit files, found in directories such as /usr/share/dbus-1/system-services/, can be replaced or added to.<sup>13</sup> This allows an adversary to hijack legitimate inter-process communication channels or inject malicious code into critical system services, often by exploiting over-permissive access policies.<sup>13</sup>

This reliance on low-level I/O events or network bindings for activation makes detection reliant not just on monitoring configuration files, but on deep kernel activity monitoring. Security tooling must correlate these low-level events (inotify/socket binding) with subsequent service activation, linking T1501 to Discovery tactics (T1518) or potential Lateral Movement tactics (T1184) if the socket activation is tied to an SSH service, for example.<sup>5</sup>

## Section 4: Threat Actors and User-Level Persistence

Real-world observations confirm that systemd persistence is a viable technique actively employed by threat groups. Furthermore, the capacity for standard users to establish persistence without root privileges broadens the attack surface significantly.

### 4.1 Case Studies of Systemd Persistence in Malware

The MITRE ATT&CK knowledge base, which is built on real-world observations, serves as a foundation for effective threat models.<sup>5</sup> The Rocke Group, a cybercrime organization known for deploying malware families written in Golang, has been documented utilizing systemd services as part of its persistence mechanisms in Linux environments.<sup>1</sup> This confirms T1501 is an

operational component of modern threat toolkits designed for backdooring Unix systems.<sup>1</sup>

## 4.2 Persistence without Root: The User-Level Threat

One of the most critical security implications of systemd is the ability for any standard, non-root user to establish reliable persistence. While system-level persistence (in `/etc/systemd/system`) requires root access, any authenticated user can bypass this requirement by deploying unit files to their dedicated user configuration directory: `$HOME/.config/systemd/user/`.<sup>1</sup>

This means a successful, non-root initial compromise—such as gaining a shell via a web vulnerability or phishing—can transition immediately into persistence without requiring a privilege escalation exploit. The adversary simply drops a malicious `.service` file in the user's home directory, schedules it using `systemctl enable --user`, and launches it with `systemctl start --user`.<sup>9</sup>

This user-level capability fundamentally changes the threat landscape. Since an attacker can institutionalize access even *before* achieving root, defenses must extend beyond monitoring system-level directories and must actively scrutinize user home directories. This persistence vector is essential for maintaining access and facilitating subsequent lateral movement or privilege escalation attempts within the organization.

# Section 5: High-Fidelity Detection and Threat Hunting Strategies

Effective detection of T1501 requires security teams to move beyond basic process monitoring and implement high-fidelity controls focused on configuration modification, administrative command execution, and low-level kernel auditing using tools such as Auditd.

## 5.1 Monitoring Critical Unit File Directories (DS0022)

Adversary activity related to systemd persistence begins with the File Creation or Modification event (DS0022) of the unit configuration files.<sup>1</sup> Robust detection engineering must prioritize monitoring specific file paths:

- **System-Wide Persistence:** High-priority monitoring should target `/etc/systemd/system/` and its subdirectories, as changes here reflect attempts at root or system-level persistence.<sup>1</sup>
- **User-Level Persistence:** Monitoring must also extend to the dynamic user configuration paths, specifically `$HOME/.config/systemd/user/`, to catch non-root persistence establishment.<sup>9</sup>

To ensure persistent monitoring across system reboots, Audit rules targeting write access (w) to these directories must be placed in `/etc/audit/rules.d/` and loaded via the `agenrules` program.<sup>8</sup> Security Information and Event Management (SIEM) solutions and Endpoint Detection and Response (EDR) tools must implement correlation rules specifically designed for the MITRE



T1501 framework.<sup>5</sup>

## 5.2 Auditing the System Control Utility (systemctl)

Persistence configuration only becomes active when systemd is instructed to load or enable the new unit file. Therefore, auditing the execution and command-line arguments of the systemctl utility, as well as related tools like `/usr/sbin/service`, is a vital detection step.<sup>1</sup>

High-priority actions to monitor include any execution of systemctl coupled with the following commands:

- `enable`: Schedules the unit for execution upon system boot.<sup>4</sup>
- `start`: Immediately executes the malicious payload.<sup>4</sup>
- `daemon-reload`: Instructs systemd to load new or modified configuration files from disk, an action necessary for activating a newly dropped unit file.<sup>9</sup>
- `link`: Used when placing symbolic links in unit directories to reference payloads elsewhere.<sup>1</sup>

The highest fidelity indicator of operational persistence deployment is the temporal correlation between a file modification event (DS0022) in a systemd unit directory and the subsequent execution of `/usr/bin/systemctl daemon-reload`. This sequence confirms that the attacker successfully wrote the malicious configuration and instructed the operating system to load and register it immediately, providing near-certain confirmation of active T1501 persistence.

## 5.3 Detection Engineering with Auditd

The foundation of robust Linux threat hunting is a properly configured Auditd system. Auditd control rules must be correctly defined, such as ensuring the system deletes existing rules on launch (`-D`) and sets the failure mode to log (`-f 1`) to maintain audit integrity.<sup>14</sup> The following table provides examples of high-priority audit rules necessary for effective T1501 detection.

Table 2: High-Priority Auditd Rules for Systemd Persistence Detection (Mapping to Audit Events)

Audit Category	Auditd Rule Type (Example)	Target Path/Argument	Purpose	Relevant ATT&CK Technique
System Unit Modification	<code>-w /etc/systemd/system -p wa -k sys_persist_mo</code>	<code>/etc/systemd/system</code> (and subdirectories)	Detect new service installation or modification of existing units. <sup>8</sup>	T1543.002: Service Modification (Root)

	d			
<b>User Unit Creation</b>	-w /home/*/*.conf systemd/user -p wa -k user_persist_add	\$HOME/.config/systemd/user/	Detect non-root persistence establishment in user contexts. <sup>9</sup>	T1543.002: Service Creation (User)
<b>Administrative Execution</b>	-a always,exit -F arch=b64 -S execve -F path=/usr/bin/systemctl -k systemctl\_exec	/usr/bin/systemctl or /usr/sbin/service	Monitor all service configuration changes, especially enable, start, and daemon-reload. <sup>1</sup>	T1501 Execution
<b>Symlink Persistence</b>	-w /etc/systemd/system -F perm=w -k symlink_abuse	/etc/systemd/system write operations	Capture placement of symbolic links used to point units to hidden payloads. <sup>1</sup>	T1501/T1108: Redundant Access

## Section 6: Mitigating T1501: Advanced Systemd Hardening (Sandboxing)

Mitigation against T1501 must incorporate proactive security measures that limit the privileges and capabilities of service processes, adhering strictly to the principle of least privilege. Systemd offers a robust set of directives designed for sandboxing, which can neutralize a malicious service even if its persistence mechanism is successfully deployed.

### 6.1 The Principle of Least Privilege in Systemd Units



Hardening constitutes a last-resort defense, focusing on mitigating the consequences of a successful exploit by limiting what a compromised process can access or execute.<sup>15</sup> The goal is to enforce a highly restrictive environment, preventing the malicious service from performing unauthorized actions, such as writing data to the filesystem, accessing the network, or escalating privileges beyond its operational requirements.<sup>15</sup>

## 6.2 Assessing and Auditing Unit Security Exposure

Systemd provides built-in mechanisms for security assessment. The utility `systemd-analyze security <unit_name>` generates a comprehensive security exposure report and a quantifiable score.<sup>16</sup> Exposure scores, such as 9.6 out of 10, are flagged as UNSAFE, signaling critical risks due to overly permissive configurations.<sup>17</sup> Hardening is therefore an iterative process, involving modifying unit directives and continuously reassessing the exposure score until the security profile meets acceptable standards.<sup>17</sup>

## 6.3 Implementing Kernel Sandboxing Directives

Systemd sandboxing relies on core kernel technologies, notably Linux Namespacing (using syscalls like `unshare`) and Seccomp filters, to isolate services.<sup>15</sup> These directives establish a defensive perimeter around the service process.

- **Filesystem Isolation:** The directive `PrivateTmp=true` isolates the service by creating a private temporary directory structure (`/tmp/systemd-private-*`) rather than utilizing the shared `/tmp` or `/var/tmp` directories.<sup>17</sup> This measure eliminates an entire class of vulnerabilities related to temporary file prediction and replacement attacks. Directives like `ReadOnlyPaths` or `ProtectSystem` can further restrict the service's write capabilities, severely limiting the attacker's ability to stage or persist data on the host filesystem.<sup>15</sup>
- **Capability and Privilege Restriction:** To cripple typical privilege escalation attempts, `NoNewPrivileges=true` must be set, preventing the service process from acquiring new privileges after its initial launch.<sup>16</sup> Furthermore, `CapabilityBoundingSet` is used to explicitly drop kernel capabilities that the service does not require. For example, dropping `CAP_SYS_TIME` prevents a compromised service from changing the system clock, an action often abused for evasion.<sup>16</sup>
- **Network and Device Isolation:** If a service does not require external communication,

setting `PrivateNetwork=true` isolates it entirely from the network stack. Similarly, `PrivateDevices` restricts access to system devices.

The fundamental challenge in mitigating T1501 is navigating the "hardening paradox," where overly strict rules risk breaking the legitimate functionality of the application.<sup>15</sup> This dictates that successful mitigation is achieved through specific tuning and runtime profiling rather than generic rules. The hardening directives act as a final, comprehensive defense layer that limits the *impact* of a successful persistence breach. If an adversary achieves T1501 persistence but the service unit is configured with `PrivateNetwork=true`, the malicious payload's ability to communicate with an external C2 infrastructure is neutralized, effectively transforming a critical security breach into a contained, non-functional persistence attempt.

## 6.4 Automated Hardening Tools

Given the complexity of manually configuring and tuning dozens of sandboxing directives, tools have emerged to automate this process. The Systemd Hardening Helper (SHH), written in Rust, is an example of such a utility. SHH profiles a service's runtime behavior and automatically generates an optimal set of strict hardening options, simplifying the necessary configuration required to secure systemd units.<sup>15</sup>

## Conclusions and Recommendations

The proliferation of systemd across modern Linux distributions has solidified T1501 (Systemd Service Persistence) as a high-priority threat vector. The analysis demonstrates that this technique is not monolithic; adversaries employ complex strategies ranging from exploiting simple `ExecStart` directives to using sophisticated, event-driven activation methods via `Timer`, `Path`, and `Socket` units. Crucially, persistence can be established without root privileges via user-level units, dramatically lowering the barrier to entry for initial access and internal persistence.

Effective defense requires a multi-layered approach:

1. **High-Fidelity Detection via Correlation:** Detection strategies must move beyond monitoring simple file changes. The highest fidelity indicators are found by correlating file write events (DS0022) in critical unit directories with the execution of administrative commands such as `systemctl daemon-reload` or `systemctl enable`.
2. **Comprehensive Directory Auditing:** Monitoring must persistently cover both

system-level directories (/etc/systemd/system) and all user-level configuration paths (\$HOME/.config/systemd/user/) to neutralize the non-root persistence vector.

3. **Proactive Sandboxing and Hardening:** Organizations should adopt a mandatory hardening posture for all services, utilizing systemd directives such as **PrivateTmp=true**, **NoNewPrivileges=true**, and **PrivateNetwork=true**. Regular security assessment using **systemd-analyze security** is necessary to continuously reduce the attack surface and limit the potential impact of a successful persistence deployment.

## Works cited

1. Create or Modify System Process: Systemd Service, Sub-technique T1543.002 - Enterprise, accessed October 29, 2025, <https://attack.mitre.org/techniques/T1543/002/>
2. Systemd Units - A Comprehensive Guide for Linux Admins - Webdock, accessed October 29, 2025, <https://webdock.io/en/docs/how-guides/system-maintenance/systemd-units-comprehensive-guide-linux-admins>
3. Defending against malware persistence techniques with Wazuh - Bleeping Computer, accessed October 29, 2025, <https://www.bleepingcomputer.com/news/security/defending-against-malware-persistence-techniques-with-wazuh/>
4. ATT&CK T1501: Understanding Systemd Service Persistence - Red Canary, accessed October 29, 2025, <https://redcanary.com/blog/threat-detection/attck-t1501-understanding-systemd-service-persistence/>
5. Linux MITRE ATTACK Rules - ScienceSoft, accessed October 29, 2025, <https://www.scnsoft.com/security/siem/linux-mitre-attack-rules>
6. Persistence, Tactic TA0003 - Enterprise | MITRE ATT&CK®, accessed October 29, 2025, <https://attack.mitre.org/tactics/TA0003/>
7. New Sub-techniques - MITRE ATT&CK®, accessed October 29, 2025, [https://attack.mitre.org/docs/Persistence\\_and\\_Privilege\\_Escalation\\_ID\\_Map\\_and\\_New\\_Techniques.xlsx](https://attack.mitre.org/docs/Persistence_and_Privilege_Escalation_ID_Map_and_New_Techniques.xlsx)
8. Chapter 37. Auditing the system | System Design Guide | Red Hat Enterprise Linux | 8, accessed October 29, 2025, [https://docs.redhat.com/en/documentation/red\\_hat\\_enterprise\\_linux/8/html/system\\_design\\_guide/auditing-the-system\\_system-design-guide](https://docs.redhat.com/en/documentation/red_hat_enterprise_linux/8/html/system_design_guide/auditing-the-system_system-design-guide)
9. Hunting for Persistence in Linux (Part 3): Systemd, Timers, and Cron - pepe berba, accessed October 29, 2025, <https://pberba.github.io/security/2022/01/30/linux-threat-hunting-for-persistence-systemd-timers-cron/>
10. Cronjob vs systemd timers - how to create a systemd timer | ryansouthgate.com, accessed October 29, 2025, <https://ryansouthgate.com/systemd-timer/>
11. Cron vs systemd timers - Unix & Linux Stack Exchange, accessed October 29, 2025, <https://unix.stackexchange.com/questions/278564/cron-vs-systemd-timers>
12. Linux Detection Engineering - The Grand Finale on Linux Persistence — Elastic

- Security Labs, accessed October 29, 2025,  
<https://www.elastic.co/security-labs/the-grand-finale-on-linux-persistence>
13. Linux detection engineering with Auditd – Elastic Security Labs, accessed October 29, 2025,  
<https://www.elastic.co/security-labs/linux-detection-engineering-with-auditd>
  14. systemd hardening made easy with SHH - Synacktiv, accessed October 29, 2025,  
<https://www.synacktiv.com/en/publications/systemd-hardening-made-easy-with-shh>
  15. Systemd Units Hardening - Rocky Linux Documentation, accessed October 29, 2025, [https://docs.rockylinux.org/10/guides/security/systemd\\_hardening/](https://docs.rockylinux.org/10/guides/security/systemd_hardening/)
  16. alegrey91/systemd-service-hardening - GitHub, accessed October 29, 2025,  
<https://github.com/alegrey91/systemd-service-hardening>