

Malloc Maleficarum

Phantasmal Phantasmagoria

October 11, 2005

Date: Tue, 11 Oct 2005 10:14:02 -0700
From: Phantasmal Phantasmagoria <phantasmal@hush.ai>
To: bugtraq@securityfocus.com
Subject: The Malloc Maleficarum
[-----

The Malloc Maleficarum
Glibc Malloc Exploitation Techniques
by Phantasmal Phantasmagoria
phantasmal@hush.ai
[-----

In late 2001, "Vudo Malloc Tricks" and "Once Upon A free()" defined the exploitation of overflowed dynamic memory chunks on Linux. In late 2004, a series of patches to GNU libc malloc implemented over a dozen mandatory integrity assertions, effectively rendering the existing techniques obsolete.

It is for this reason, a small suggestion of impossibility, that I present the Malloc Maleficarum.

[-----
The House of Prime
The House of Mind
The House of Force
The House of Lore
The House of Spirit
The House of Chaos
[-----

1 The House of Prime

An artist has the first brush stroke of a painting. A writer has the first line of a poem. I have the House of Prime. It was the first breakthrough, the indication of everything that was to come. It was the rejection of impossibility. And it was also the most difficult to derive. For these reasons I feel obliged to give Prime the position it deserves as the first House of the Malloc Maleficarum.

>From a purely technical perspective the House of Prime is perhaps the least useful of the collection. It is almost invariably better to use the House of Mind or Spirit when the conditions allow it. In order to successfully apply the House of Prime it must be possible to free() two different chunks with designer controlled size fields and then trigger a call to malloc().

The general idea of the technique is to corrupt the fastbin maximum size variable, which under certain uncontrollable circumstances (discussed below) allows the designer to hijack the arena structure used by calls to `malloc()`, which in turn allows either the return of an arbitrary memory chunk, or the direct modification of execution control data.

As previously stated, the technique starts with a call to `free()` on an area of memory that is under control of the designer. A call to `free()` actually invokes a wrapper, called `public_fREe()`, to the internal function `_int_free()`. For the House of Prime, the details of `public_fREe()` are relatively unimportant. So attention moves, instead, to `_int_free()`. From the glibc-2.3.5 source code:

```

1
2 void
3 _int_free(mstate av, Void_t* mem)
4 {
5     mchunkptr      p;           /* chunk corresponding to mem */
6     INTERNAL_SIZE_T size;       /* its size */
7     mfastbinptr*    fb;         /* associated fastbin */
8     ...
9
10    p = mem2chunk(mem);
11    size = chunksize(p);
12
13    if (__builtin_expect ((uintptr_t) p > (uintptr_t) -size, 0)
14        || __builtin_expect ((uintptr_t) p & MALLOC_ALIGN_MASK, 0))
15    {
16        errstr = "free(): invalid pointer";
17        errout:
18        malloc_printerr (check_action, errstr, mem);
19        return;
20    }

```

Almost immediately one of the much vaunted integrity tests appears. The `__builtin_expect()` construct is used for optimization purposes, and does not in any way effect the conditions it contains. The designer must ensure that both of the tests fail in order to continue execution. At this stage, however, doing so is not difficult.

Note that the designer does not control the value of `p`. It can therefore be assumed that the test for misalignment will fail. On the other hand, the designer does control the value of `size`. In fact, it is the most important aspect of control that the designer possesses, yet its range is already being limited. For the the House of Prime the exact upper limit of `size` is not important. The lower limit, however, is crucial in the correct execution of this technique. The `chunksize()` macro is defined as follows:

```

1
2 #define SIZE_BITS (PREV_INUSE|IS_MMAPPED|NON_MAIN_ARENA)
3 #define chunksize(p) ((p)->size & ~(SIZE_BITS))

```

The `PREV_INUSE`, `IS_MMAPPED` and `NON_MAIN_ARENA` definitions correspond to the three least significant bits of the `size` entry in a `malloc` chunk. The `chunksize()` macro clears these three bits, meaning the lowest possible value of the designer controlled `size` value is 8. Continuing with `_int_free()` it will soon become clear why this is important:

```

1
2 if ((unsigned long)(size) <= (unsigned long)(av->max_fast))
3 {
4     if (chunk_at_offset (p, size)->size <= 2 * SIZE_SZ

```

```

5         || __builtin_expect (chunksize (chunk_at_offset (p, size))
6                               >= av->$system_mem, 0))
7     {
8         errstr = "free(): invalid next size (fast)";
9         goto errout;
10    }
11
12    set_fastchunks(av);
13    fb = &(av->$fastbins[fastbin_index(size)]);
14
15    if (__builtin_expect (*fb == p, 0))
16    {
17        errstr = "double free or corruption (fasttop)";
18        goto errout;
19    }
20
21    p->fd = *fb;
22    *fb = p;
23 }

```

This is the fastbin code. Exactly what a fastbin is and why they are used is beyond the scope of this document, but remember that the first step in the House of Prime is to overwrite the fastbin maximum size variable, `av->max_fast`. In order to do this the designer must first provide a chunk with the lower limit size, which was derived above. Given that the default value of `av->max_fast` is 72 it is clear that the fastbin code will be used for such a small size. However, exactly why this results in the corruption of `av->max_fast` is not immediately apparent.

It should be mentioned that `av` is the arena pointer. The arena is a control structure that contains, amongst other things, the maximum size of a fastbin and an array of pointers to the fastbins themselves. In fact, `av->max_fast` and `av->fastbins` are contiguous:

```

1
2    ...
3    INTERNAL_SIZE_T    max_fast;
4    mfastbinptr        fastbins[NFASTBINS];
5    mchunkptr          top;
6    ...

```

Assuming that the nextsize integrity check fails, the `fb` pointer is set to the address of the relevant fastbin for the given size. This is computed as an index from the zeroth entry of `av->fastbins`. The zeroth entry, however, is designed to hold chunks of a minimum size of 16 (the minimum size of a malloc chunk including `prev_size` and `size` values). So what happens when the designer supplies the lower limit size of 8? An analysis of `fastbin_index()` is needed:

```

1
2 #define fastbin_index(sz)      (((unsigned int)(sz)) >> 3) - 2)

```

Simple arithmetic shows that $8 \gg 3 = 1$, and $1 - 2 = -1$. Therefore `fastbin_index(8)` is -1, and thus `fb` is set to the address of `av->fastbins[-1]`. Since `av->max_fast` is contiguous to `av->fastbins` it is evident that the `fb` pointer is set to `&av->max_fast`. Furthermore, the second integrity test fails (since `fb` definitely does not point to `p`) and the final two lines of the fastbin code are reached. Thus the forward pointer of the designer's chunk `p` is set to `av->max_fast`, and `av->max_fast` is set to the value of `p`.

An assumption was made above that the nextsize integrity check fails. In reality it often takes a bit of work to get this to fall together. If the overflow is capable of writing null bytes, then the solution is simple.

However, if the overflow terminates on a null byte, then the solution becomes application specific. If the tests fail because of the natural memory layout at overflow, which they often will, then there is no problem. Otherwise some memory layout manipulation may be needed to ensure that the nextsize value is designer controlled.

The challenging part of the House of Prime, however, is not how to overwrite `av->max_fast`, but how to leverage the overwrite into arbitrary code execution. The House of Prime does this by overwriting a thread specific data variable called `arena_key`. This is where the biggest condition of the House of Prime arises. Firstly, `arena_key` only exists if glibc malloc is compiled with `USE_ARENAS` defined (this is the default setting). Furthermore, and most significantly, `arena_key` must be at a higher address than the actual arena:

```
0xb7f00000 <main_arena>:      0x00000000
0xb7f00004 <main_arena+4>:    0x00000049      <-- max_fast
0xb7f00008 <main_arena+8>:    0x00000000      <-- fastbin[0]
0xb7f0000c <main_arena+12>:   0x00000000      <-- fastbin[1]
....
0xb7f00488 <mp_+40>:          0x0804a000      <-- mp_.sbrk_base
0xb7f0048c <arena_key>:      0xb7f00000
```

Due to the fact that the arena structure and the `arena_key` come from different source files, exactly when this does and doesn't happen depends on how the target libc was compiled and linked. I have seen the cards fall both ways, so it is an important point to make. For now it will be assumed that the `arena_key` is at a higher address, and is thus over-writable by the fastbin code.

The `arena_key` is thread specific data, which simply means that every thread of execution has its own `arena_key` independent of other threads. This may have to be considered when applying the House of Prime to a threaded program, but otherwise `arena_key` can safely be treated as normal data.

The `arena_key` is an interesting target because it is used by the `arena_get()` macro to find the arena for the currently executing thread. That is, if `arena_key` is controlled for some thread and a call to `arena_get()` is made, then the arena can be hijacked. Arena hijacking of this type will be covered shortly, but first the actual overwrite of `arena_key` must be considered.

In order to overwrite `arena_key` the fastbin code is used for a second time. This corresponds to the second `free()` of a designer controlled chunk that was outlined in the original prerequisites for the House of Prime. Normally the fastbin code would not be able to write beyond the end of `av->fastbins`, but since `av->max_fast` has previously been corrupted, chunks with any size less than the value of the address of the designer's first chunk will be treated with the fastbin code. Thus the designer can write up to `av->fastbins[fastbin_index(av->max_fast)]`, which is easily a large enough range to be able to reach the `arena_key`.

In the example memory dump provided above the `arena_key` is 0x484 (1156) bytes from `av->fastbins[0]`. Therefore an index of `1156/sizeof(mfastbinptr)` is needed to set `fb` to the address of `arena_key`. Assuming that the system has 32-bit pointers a `fastbin_index()` of 289 is required. Roughly inverting the `fastbin_index()` gives:

$$(289 + 2) \ll 3 = 2328$$

This means that a size of 2328 will result in `fb` being set to `arena_key`. Note that this size only applies for the memory dump shown above. It is quite likely that the offset between `av->fastbins[0]` and `arena_key`

will differ from system to system.

Now, if the designer has corrupted `av->max_fast` and triggered a `free()` on a chunk with size 2328, and assuming the failure of the nextsize integrity tests, then `fb` will be set to `arena_key`, the forward pointer of the designer's second chunk will be set to the address of the existing arena, and `arena_key` will be set to the address of the designer's second chunk.

When corrupting `av->max_fast` it was not important for the designer to control the overflowed chunk so long as the nextsize integrity checks were handled. When overwriting `arena_key`, however, it is crucial that the designer controls at least part of the overflowed chunk's data. This is because the overflowed chunk will soon become the new arena, so it is natural that at least part of the chunk data must be arbitrarily controlled, or else arbitrary control of the result of `malloc()` could not be expected.

A call to `malloc()` invokes a wrapper function called `public_mALLOc()`:

```
1
2 Void_t*
3 public_mALLOc(size_t bytes)
4 {
5     mstate ar_ptr;
6     Void_t *victim;
7     ...
8     arena_get(ar_ptr, bytes);
9     if(!ar_ptr)
10         return 0;
11     victim = _int_malloc(ar_ptr, bytes);
12     ...
13     return victim;
14 }
```

The `arena_get()` macro is in charge of finding the current arena by retrieving the `arena_key` thread specific data, or failing this, creating a new arena. Since the `arena_key` has been overwritten with a non-zero quantity it can be safely assumed that `arena_get()` will not try to create a new arena. In the `public_mALLOc()` wrapper this has the effect of setting `ar_ptr` to the new value of `arena_key`, the address of the designer's second chunk. In turn this value is passed to the internal function `_int_malloc()` along with the requested allocation size.

Once execution passes to `_int_malloc()` there are two ways for the designer to proceed. The first is to use the fastbin allocation code:

```
1
2 Void_t*
3 _int_malloc(mstate av, size_t bytes)
4 {
5     INTERNAL_SIZE_T nb;           /* normalized request size */
6     unsigned int idx;             /* associated bin index */
7     mfastbinptr* fb;              /* associated fastbin */
8     mchunkptr victim;             /* inspected/selected chunk */
9
10    checked_request2size(bytes, nb);
11
12    if ((unsigned long)(nb) <= (unsigned long)(av->max_fast)) {
13        long int idx = fastbin_index(nb);
14        fb = &(av->fastbins[idx]);
```

```

15     if ( (victim = *fb) != 0) {
16         if (fastbin_index (chunksize (victim)) != idx)
17             malloc_printerr (check_action, "malloc(): memory"
18                 " corruption (fast)", chunk2mem (victim));
19         *fb = victim$->$fd;
20         check_reallocated_chunk(av, victim, nb);
21         return chunk2mem(victim);
22     }
23 }

```

The `checked_request2size()` macro simply converts the request into the absolute size of a memory chunk with data length of the requested size. Remember that `av` is pointing towards a designer controlled area of memory, and also that the forward pointer of this chunk has been corrupted by the fastbin code. If glibc malloc is compiled without thread statistics (which is the default), then `p->fd` of the designer's chunk corresponds to `av->fastbins[0]` of the designer's arena. For the purposes of this technique the use of `av->fastbins[0]` must be avoided. This means that the request size must be greater than 8.

Interestingly enough, if the absence of thread statistics is assumed, then `av->max_fast` corresponds to `p->size`. This has the effect of forcing `nb` to be less than the size of the designer's second chunk, which in the example provided was 2328. If this is not possible, the designer must use the `unsorted_chunks/largebin` technique that will be discussed shortly.

By setting up a fake fastbin entry at `av->fastbins[fastbin_index(nb)]` it is possible to return a chunk of memory that is actually on the stack. In order to pass the `victimsize` integrity test it is necessary to point the fake fastbin at a user controlled value. Specifically, the size of the victim chunk must have the same `fastbin_index()` as `nb`, so the fake fastbin must point to 4 bytes before the designer's value in order to have the right positioning for the call to `chunksize()`.

Assuming that there is a designer controlled variable on the stack, the application will subsequently handle the returned area as if it were a normal memory chunk of the requested size. So if there is a saved return address in the "allocated" range, and if the designer can control what the application writes to this range, then it is possible to circumvent execution to an arbitrary location.

If it is possible to trigger an appropriate `malloc()` with a request size greater than the size of the designer's second chunk, then it is better to use the `unsorted_chunks` code in `_int_malloc()` to cause an arbitrary memory overwrite. This technique does, however, require a greater amount of designer control in the second chunk, and further control of two areas of memory somewhere in the target process address space. To trigger the `unsorted_chunks` code at all the absolute request size must be larger than 512 (the maximum smallbin chunk size), and of course, must be greater than the fake arena's `av->max_fast`. Assuming it is, the `unsorted_chunks` code is reached:

```

1     for(;;) {
2         while ( (victim = unsorted_chunks(av)->bk) !=
3 unsorted_chunks(av)) {
4             bck = victim$->$bk;
5
6
7             if (__builtin_expect (victim$->$size <= 2 * SIZE_SZ, 0)
8                 || __builtin_expect (victim$->$size > av$->$system_mem, 0))
9                 malloc_printerr (check_action, "malloc(): memory"
10                    " corruption", chunk2mem (victim));
11

```

```

12     size = chunksize(victim);
13
14     if (in_smallbin_range(nb) &&
15         bck == unsorted_chunks(av) &&
16         victim == av->$last_remainder &&
17         (unsigned long)(size) > (unsigned long)(nb + MINSIZE)) {
18         ...
19     }
20
21     unsorted_chunks(av)->bk = bck;
22     bck->fd = unsorted_chunks(av);
23
24     if (size == nb) {
25         ...
26         return chunk2mem(victim);
27     }
28     ...

```

There are quite a lot of things to consider here. Firstly, the `unsorted_chunks()` macro returns `av->bins[0]`. Since the designer controls `av`, the designer also controls the value of `unsorted_chunks()`. This means that `victim` can be set to an arbitrary address by creating a fake `av->bins[0]` value that points to an area of memory (called A) that is designer controlled. In turn, `A->bk` will contain the address that `victim` will be set to (called B). Since `victim` is at an arbitrary address B that can be designer controlled, the temporary variable `bck` can be set to an arbitrary address from `B->bk`.

For the purposes of this technique, `B->size` should be equal to `nb`. This is not strictly necessary, but works well to pass the two `victimsize` integrity tests while also triggering the final condition shown above, which has the effect of ending the call to `malloc()`.

Since it is possible to set `bck` to an arbitrary location, and since `unsorted_chunks()` returns the designer controlled area of memory A, the setting of `bck->fd` to `unsorted_chunks()` makes it possible to set any location in the address space to A. Redirecting execution is then a simple matter of setting `bck` to the address of a GOT or `..dtors` entry minus 8. This will redirect execution to `A->prev_size`, which can safely contain a near `jmp` to skip past the crafted value at `A->bk`. Similar to the fastbin allocation code the arbitrary address B is returned to the requesting application.

2 The House of Mind

Perhaps the most useful and certainly the most general technique in the Malloc Maleficarum is the House of Mind. The House of Mind has the distinct advantage of causing a direct memory overwrite with just a single call to `free()`. In this sense it is the closest relative in the Malloc Maleficarum to the traditional `unlink()` technique.

The method used involves tricking the wrapper invoked by `free()`, called `public_fREe()`, into supplying the `_int_free()` internal function with a designer controlled arena. This can subsequently lead to an arbitrary memory overwrite. A call to `free()` actually invokes a wrapper called `public_fREe()`:

```

1
2 void
3 public_fREe(Void_t* mem)
4 {

```

```

5     mstate ar_ptr;
6     mchunkptr p;          /* chunk corresponding to mem */
7     ...
8     p = mem2chunk(mem);
9     ...
10    ar_ptr = arena_for_chunk(p);
11    ...
12    _int_free(ar_ptr, mem);

```

When memory is passed to `free()` it points to the start of the data portion of the "corresponding chunk". In an allocated state a chunk consists of the `prev_size` and `size` values and then the data section itself. The `mem2chunk()` macro is in charge of converting the supplied memory value into the corresponding chunk. This chunk is then passed to the `arena_for_chunk()` macro:

```

1
2 #define HEAP_MAX_SIZE (1024*1024) /* must be a power of two */
3
4 #define heap_for_ptr(ptr) \
5     ((heap_info *)((unsigned long)(ptr) & ~(HEAP_MAX_SIZE-1)))
6
7 #define chunk_NON_MAIN_ARENA(p) ((p)->size & NON_MAIN_ARENA)
8
9 #define arena_for_chunk(ptr) \
10    (chunk_NON_MAIN_ARENA(ptr)?heap_for_ptr(ptr)->ar_ptr:&main_arena)

```

The `arena_for_chunk()` macro is tasked with finding the appropriate arena for the chunk in question. If glibc malloc is compiled with `USE_ARENAS` (which is the default), then the code shown above is used. Clearly, if the `NON_MAIN_ARENA` bit in the size value of the chunk is not set, then `ar_ptr` will be set to the `main_arena`.

However, since the designer controls the size value it is possible to control whether the chunk is treated as being in the main arena or not. This is what the `chunk_NON_MAIN_ARENA()` macro checks for. If the `NON_MAIN_ARENA` bit is set, then `chunk_NON_MAIN_ARENA()` returns positive and `ar_ptr` is set to `heap_for_ptr(ptr)->ar_ptr`.

When a non-main heap is created it is aligned to a multiple of `HEAP_MAX_SIZE`. The first thing that goes into this heap is the `heap_info` structure. Most significantly, this structure contains an element called `ar_ptr`, the pointer to the arena for this heap. This is how the `heap_for_ptr()` macro functions, aligning the given chunk down to a multiple of `HEAP_MAX_SIZE` and taking the `ar_ptr` from the resulting `heap_info` structure.

The House of Mind works by manipulating the heap so that the designer controls the area of memory that the overflowed chunk is aligned down to. If this can be achieved, an arbitrary `ar_ptr` value can be supplied to `_int_free()` and subsequently an arbitrary memory overwrite can be triggered. Manipulating the heap generally involves forcing the application to repeatedly allocate memory until a designer controlled buffer is contained at a `HEAP_MAX_SIZE` boundary.

In practice this alignment is necessary because chunks at low areas of the heap align down to an area of memory that is neither designer controlled nor mapped in to the address space. Fortunately, the amount of allocation that creates the correct alignment is not large. With the default `HEAP_MAX_SIZE` of `1024*1024` an average of 512kb of padding will be required, with this figure never exceeding 1 megabyte.

It should be noted that there is not a general method for triggering memory allocation as required by the House of Mind, rather the process is application specific. If a situation arises in which it is impossible to align a designer controlled chunk, then the House of Lore or Spirit should be considered.

So, it is possible to hijack the `heap_info` structure used by the `heap_for_ptr()` macro, and thus supply an arbitrary value for `ar_ptr` which controls the arena used by `_int_free()`. At this stage the next question that arises is exactly what to do with `ar_ptr`. There are two options, each with their respective advantages and disadvantages. Each will be addressed in turn.

Firstly, setting the `ar_ptr` to a sufficiently large area of memory that is under the control of the designer and subsequently using the unsorted chunk link code to cause a memory overwrite. Sufficiently large in this case means the size of the arena structure, which is 1856 bytes on a 32-bit system without `THREAD_STATS` enabled. The main difficulty in this method arises with the numerous integrity checks that are encountered. Fortunately, nearly every one of these tests use a value obtained from the designer controlled arena, which makes the checks considerably easier to manage.

For the sake of brevity, the complete excerpt leading up to the unsorted chunk link code has been omitted. Instead, the following list of the conditions required to reach the code in question is provided. Note that both `av` and the size of the overflowed chunk are designer controlled values.

- The negative of the size of the overflowed chunk must be less than the value of the chunk itself.
- The size of the chunk must not be less than `av->max_fast`.
- The `IS_MMAPPED` bit of the size cannot be set.
- The overflowed chunk cannot equal `av->top`.
- The `NONCONTIGUOUS_BIT` of `av->max_fast` must be set.
- The `PREV_INUSE` bit of the nextchunk (`chunk + size`) must be set.
- The size of nextchunk must be greater than 8.
- The size of nextchunk must be less than `av->system_mem`
- The `PREV_INUSE` bit of the chunk must not be set.
- The nextchunk cannot equal `av->top`.
- The `PREV_INUSE` bit of the chunk after nextchunk (`nextchunk + nextsize`) must be set

If these conditions are met, then the following code is reached:

```
1      bck = unsorted_chunks(av);  
2      fwd = bck->fd;  
3
```

```

4     p->bk = bck;
5     p->fd = fwd;
6     bck->fd = p;
7     fwd->bk = p;

```

In this case `p` is the address of the designer's overflowed chunk. The `unsorted_chunks()` macro returns `av->bins[0]` which is designer controlled. If the designer sets `av->bins[0]` to the address of a GOT or `.dtors` entry minus 8, then that entry (`bck->fd`) will be overwritten with the address of `p`. This address corresponds to the `prev_size` entry of the designer's overflowed chunk which can safely be used to branch past the corrupted size, `fd` and `bk` entries.

The extensive list of conditions appear to make this method quite difficult to apply. In reality, the only conditions that may be a problem are those involving the `nextchunk`. This is because they largely depend on the application specific memory layout to handle. This is the only obvious disadvantage of the method. As it stands, the House of Mind is in a far better position than the House of Prime to handle such conditions due to the arbitrary nature of `av->system_mem`.

It should be noted that the last element of the arena structure that is actually required to reach the unsorted chunk link code is `av->system_mem`, but it is not terribly important what this value is so long as it is high. Thus if the conditions are right, it may be possible to use this method with only 312 bytes of designer controlled memory. However, even if there is not enough designer controlled memory for this method, the House of Mind may still be possible with the second method.

The second method uses the fastbin code to cause a memory overwrite. The main advantage of this method is that it is not necessary to point `ar_ptr` at designer controlled memory, and that there are considerably less integrity checks to worry about. Consider the fastbin code:

```

1
2     if ((unsigned long)(size) <= (unsigned long)(av->$max_fast))
3     {
4         if (chunk_at_offset (p, size)->size <= 2 * SIZE_SZ
5             || __builtin_expect (chunksize (chunk_at_offset (p, size))
6                                 >= av->$system_mem, 0))
7         {
8             errstr = "free(): invalid next size (fast)";
9             goto errout;
10        }
11
12        set_fastchunks(av);
13        fb = &(av->$fastbins[fastbin_index(size)]);
14        ...
15        p->fd = *fb;
16        *fb = p;
17    }

```

The ultimate goal here is to set `fb` to the address of a GOT or `.dtors` entry, which subsequently gets set to the address of the designer's overflowed chunk. However, in order to reach the final line a number of conditions must still be met. Firstly, `av->max_fast` must be large enough to trigger the fastbin code at all. Then the size of the `nextchunk` (`p + size`) must be greater than 8, while also being less than `av->system_mem`.

The tricky part of this method is positioning `ar_ptr` in a way such that both the `av->max_fast` element at `(av + 4)` and the `av->system_mem` element at `(av + 1848)` are large enough. If a binary

has a particularly small GOT table, then it is quite possible that the highest available large number for `av->system_mem` will result in an `av->max_fast` that is actually in the area of unmapped memory between the text and data segments. In practice this shouldn't occur very often, and if it does, then the stack may be used to a similar effect.

For more information on the fastbin code, including a description of `fastbin_index()` that will help in positioning `fb` to a GOT or `..dtors` entry, consult the House of Prime.

3 The House of Force

I first wrote about glibc malloc in 2004 with "Exploiting the Wilderness". Since the techniques developed in that text were some of the first to become obsolete, and since the Malloc Maleficarum was written in the spirit of continuation and progress, I feel obliged to include another attempt at exploiting the wilderness. This is the purpose of the House of Force. From "Exploiting the Wilderness":

"The wilderness is the top-most chunk in allocated memory. It is similar to any normal malloc chunk - it has a chunk header followed by a variably long data section. The important difference lies in the fact that the wilderness, also called the top chunk, borders the end of available memory and is the only chunk that can be extended or shortened. This means it must be treated specially to ensure it always exists; it must be preserved."

So the glibc malloc implementation treats the wilderness as a special case in calls to `malloc()`. Furthermore, the top chunk will realistically never be passed to a call to `free()` and will never contain application data. This means that if the designer can trigger a condition that only ever results in the overflow of the top chunk, then the House of Force is the only option (in the Malloc Maleficarum at least).

The House of Force works by tricking the top code in to setting the wilderness pointer to an arbitrary value, which can result in an arbitrary chunk of data being returned to the requesting application. This requires two calls to `malloc()`. The major disadvantage of the House of Force is that the first call must have a completely designer controlled request size. The second call must simply be large enough to trigger the wilderness code, while the chunk returned must be (to some extent) designer controlled.

The following is the wilderness code with some additional context:

```
1
2 Void_t*
3 _int_malloc(mstate av, size_t bytes)
4 {
5     INTERNAL_SIZE_T nb;           /* normalized request size */
6     mchunkptr victim;             /* inspected/selected chunk */
7     INTERNAL_SIZE_T size;         /* its size */
8     mchunkptr remainder;         /* remainder from a split */
9     unsigned long remainder_size; /* its size */
10    ...
11    checked_request2size(bytes, nb);
12    ...
13    use_top:
14        victim = av->$top;
15        size = chunksize(victim);
16
```

```

17     if ((unsigned long)(size) >= (unsigned long)(nb + MINSIZE)) {
18         remainder_size = size - nb;
19         remainder = chunk_at_offset(victim, nb);
20         av$->$top = remainder;
21         set_head(victim, nb | PREV_INUSE |
22                 (av != &main_arena ? NON_MAIN_ARENA : 0));
23         set_head(remainder, remainder_size | PREV_INUSE);
24         check_malloced_chunk(av, victim, nb);
25         return chunk2mem(victim);
26     }

```

The first goal of the House of Force is to overwrite the wilderness pointer, `av->top`, with an arbitrary value. In order to do this the designer must have control of the location of the remainder chunk. Assume that the existing top chunk has been overflowed resulting in the largest possible size (preferably `0xffffffff`). This is done to ensure that even large values passed as an argument to `malloc` will trigger the wilderness code instead of trying to extend the heap.

The `checked_request2size()` macro ensures that the requested value is less than `-2*MINSIZE` (by default `-32`), while also adding on enough room for the `size` and `prev_size` fields and storing the final value in `nb`. For the purposes of this technique the `checked_request2size()` macro is relatively unimportant.

It was previously mentioned that the first call to `malloc()` in the House of Force must have a designer controlled argument. It can be seen that the value of `remainder` is obtained by adding the request size to the existing top chunk. Since the top chunk is not yet under the designer's control the request size must be used to position `remainder` to at least 8 bytes before a `.GOT` or `.dtors` entry, or any other area of memory that may subsequently be used by the designer to circumvent execution.

Once the wilderness pointer has been set to the arbitrary remainder chunk, any calls to `malloc()` with a large enough request size to trigger the top chunk will be serviced by the designer's wilderness. Thus the only restriction on the new wilderness is that the size must be larger than the request that is triggering the top code. In the case of the wilderness being set to overflow a `GOT` entry this is never a problem. It is then simply a matter of finding an application specific scenario in which such a call to `malloc()` is used for a designer controlled buffer.

The most important issue concerning the House of Force is exactly how to get complete control of the argument passed to `malloc()`. Certainly, it is extremely common to have at least some degree of control over this value, but in order to complete the House of Force, the designer must supply an extremely large and specifically crafted value. Thus it is unlikely to get a sufficient value out of a situation like:

```

1     buf = (char *) malloc(strlen(str) + 1);
2

```

Rather, an acceptable scenario is much more likely to be an integer variable passed as an argument to `malloc()` where the variable has previously been set by, for example, a designer controlled `read()` or `atoi()`.

4 The House of Lore

The House of Lore came to me as I was reviewing the draft write-up of the House of Prime. When I first derived the House of Prime my main concern was how to leverage the particularly overwrite that a high `av->max_fast` in the fastbin code allowed. Upon reconsideration of the problem I realized that in my first

take of the potential overwrite targets I had completely overlooked the possibility of corrupting a bin entry.

As it turns out, it is not possible to leverage a corrupted bin entry in the House of Prime since `av->max_fast` is large and the bin code is never executed. However, during this process of elimination I realized that if a bin were to be corrupted when `av->max_fast` was not large, then it might be possible to control the return value of a `malloc()` request.

At this stage I began to consider the application of bin corruption to a general malloc chunk overflow. The question was whether a linear overflow of a malloc chunk could result in the corruption of a bin. It turns out that the answer to this is, quite simply, yes it could. Furthermore, if the designer's ability to manipulate the heap is limited, or if none of the other Houses can be applied, then bin corruption of this type can in fact be very useful.

The House of Lore works by corrupting a bin entry, which can subsequently lead to `malloc()` returning an arbitrary chunk. Two methods of bin corruption are presented here, corresponding to the overflow of both small and large bin entries. The general method involves overwriting the linked list data of a chunk previously processed by `free()`. In this sense the House of Lore is quite similar to the `frontlink()` technique presented in "Vudo Malloc Tricks".

The conditions surrounding the House of Lore are quite unique. Fundamentally, the method targets a chunk that has already been processed by `free()`. Because of this it is reasonable to assume that the chunk will not be passed to `free()` again. This means that in order to leverage such an overflow only calls to `malloc()` can be used, a property shared only by the House of Force. The first method will use the `smallbin` allocation code:

```
1
2 Void_t*
3 _int_malloc(mstate av, size_t bytes)
4 {
5     ....
6     checked_request2size(bytes, nb);
7
8     if ((unsigned long)(nb) <= (unsigned long)(av->max_fast)) {
9         ...
10    }
11
12    if (in_smallbin_range(nb)) {
13        idx = smallbin_index(nb);
14        bin = bin_at(av, idx);
15
16        if ( (victim = last(bin)) != bin) {
17            if (victim == 0) /* initialization check */
18                malloc_consolidate(av);
19            else {
20                bck = victim->bk;
21                set_inuse_bit_at_offset(victim, nb);
22                bin->bk = bck;
23                bck->fd = bin;
24                ...
25                return chunk2mem(victim);
26            }
27        }
28    }
```

So, assuming that a call to `malloc()` requests more than `av->max_fast` (default 72) bytes, the check for a "smallbin" chunk is reached. The `in_smallbin_range()` macro simply checks that the request is less than the maximum size of a smallbin chunk, which is 512 by default. The smallbins are unique in the sense that there is a bin for every possible chunk size between `av->max_fast` and the smallbin maximum. This means that for any given `smallbin_index()` the resulting bin, if not empty, will contain a chunk to fit the request size.

It should be noted that when a chunk is passed to `free()` it does not go directly in to its respective bin. It is first put on the "unsorted chunk" bin. If the next call to `malloc()` cannot be serviced by an existing smallbin chunk or the unsorted chunk itself, then the unsorted chunks are sorted in to the appropriate bins. For the purposes of the House of Lore, overflowing an unsorted chunk is not very useful. It is necessary then to ensure that the chunk being overflowed has previously been sorted into a bin by `malloc()`.

Note that in order to reach the actual smallbin unlink code there must be at least one chunk in the bin corresponding to the `smallbin_index()` for the current request. Assume that a small chunk of data size N has previously been passed to `free()`, and that it has made its way into the corresponding smallbin for chunks of absolute size $(N + 8)$. Assume that the designer can overflow this chunk with arbitrary data. Assume also that the designer can subsequently trigger a call to `malloc()` with a request size of N .

If all of this is possible, then the smallbin unlink code can be reached. When a chunk is removed from the unsorted bin it is put at the front of its respective small or large bin. When a chunk is taken off a bin, such as during the smallbin unlink code, it is taken from the end of the bin. This is what the `last()` macro does, find the last entry in the requested bin. So, effectively the "victim" chunk in the smallbin unlink code is taken from `bin->bk`. This means that in order to reach the designer's victim chunk it may be necessary to repeat the N sized `malloc()` a number of times.

It should be stressed that the goal of the House of Lore to control the `bin->bk` value, but at this stage only `victim->bk` is controlled. So, assuming that the designer can trigger a `malloc()` that results in an overflowed victim chunk being passed to the smallbin unlink code, the designer (as a result of the control of `victim->bk`) controls the value of `bck`. Since `bin->bk` is subsequently set to `bck`, `bin->bk` can be arbitrarily controlled. The only condition to this is that `bck` must point to an area of writable memory due to `bck_jfd` being set at the final stage of the unlinking process.

The question then lies in how to leverage this smallbin corruption. Since the `malloc()` call that the designer used to gain control of `bin->bk` immediately returns the victim chunk to the application, at least one more call to `malloc()` with the same request size N is needed. Since `bin->bk` is under the designer's control so is `last(bin)`, and thus so is `victim`. The only thing preventing an arbitrary victim chunk being returned to the application is the fact that `bck`, set from `victim->bck`, must point to writable memory.

This rules out pointing the victim chunk at a GOT or `.dtors` entry. Instead, the designer must point `victim` to a position on the stack such that `victim->bk` is a pointer to writable memory yet still close enough to a saved return address such that it can be overwritten by the application's general use of the chunk. Alternatively, an application specific approach may be taken that targets the use of function pointers. Whichever method used, the arbitrary `malloc()` chunk must be designer controlled to some extent during its use by the application.

For the House of Lore, the only other interesting situation is when the overflowed chunk is large. In this context large means anything bigger than the maximum smallbin chunk size. Again, it is necessary for

the overflowed chunk to have previously been processed by `free()` and to have been put into a largebin by `malloc()`.

The general method of using largebin corruption to return an arbitrary chunk is similar to the case of a smallbin in the sense that the initial bin corruption occurs when an overflowed victim chunk is handled by the largebin unlink code, and that a subsequent large request will use the corrupted bin to return an arbitrary chunk. However, the largebin code is significantly more complex in comparison. This means that the conditions required to cause and leverage a bin corruption are slightly more restrictive.

The entire largebin implementation is much too large to present in full, so a description of the conditions that cause the largebin unlink code to be executed will have to suffice. If the designer's overflowed chunk of size N is in a largebin, then a subsequent request to allocate N bytes will trigger a block of code that searches the corresponding bin for an available chunk, which will eventually find the chunk that was overflowed. However, this particular block of code uses the `unlink()` macro to remove the designer's chunk from the bin. Since the `unlink()` macro is no longer an interesting target, this situation must be avoided.

So in order to corrupt a largebin a request to allocate M bytes is made, such that $512 < M < N$. If there are no appropriate chunks in the bin corresponding to requests of size M , then glibc malloc iterates through the bins until a sufficiently large chunk is found. If such a chunk is found, then the following code is used:

```
1      victim = last(bin);
2      ..
3      size = chunksize(victim);
4      remainder_size = size - nb;
5
6      bck = victim->$bk;
7      bin->bk = bck;
8      bck->fd = bin;
9
10
11     if (remainder_size < MINSIZE) {
12         set_inuse_bit_at_offset(victim, size);
13         ...
14         return chunk2mem(victim);
15     }
```

If the victim chunk is the designer's overflowed chunk, then the situation is almost exactly equivalent to the smallbin unlink code. If the designer can trigger enough calls to `malloc()` with a request of M bytes so that the overflowed chunk is used here, then the `bin->bk` value can be set to an arbitrary value and any subsequent call to `malloc()` of size Q ($512 < Q < N$) that tries to allocate a chunk from the bin that has been corrupted will result in an arbitrary chunk being returned to the application.

There are only two conditions. The first is exactly the same as the case of smallbin corruption, the `bk` pointer of the arbitrary chunk being returned to the application must point to writable memory (or the setting of `bck->fd` will cause a segmentation fault).

The other condition is not obvious from the limited code that has been presented above. If the `remainder_size` value is not less than `MINSIZE`, then glibc malloc attempts to split off a chunk at `victim + nb`. This includes calling the `set_foot()` macro with `victim + nb` and `remainder_size` as arguments. In effect, this tries to set `victim + nb + remainder_size` to `remainder_size`. If the `chunksize(victim)` (and thus

remainder_size) is not designer controlled, then set_foot() will likely try to set an area of memory that isn't mapped in to the address space (or is read-only).

So, in order to prevent set_foot() from crashing the process the designer must control both victim->size and victim->bk of the arbitrary victim chunk that will be returned to the application. If this is possible, then it is advisable to trigger the condition shown in the code above by forcing remainder_size to be less than MINSIZE. This is recommended because the condition minimizes the amount of general corruption caused, simply setting the inuse bit at victim + size and then returning the arbitrary chunk as desired.

5 The House of Spirit

The House of Spirit is primarily interesting because of the nature of the circumstances leading to its application. It is the only House in the Malloc Maleficarum that can be used to leverage both a heap and stack overflow. This is because the first step is not to control the header information of a chunk, but to control a pointer that is passed to free(). Whether this pointer is on the heap or not is largely irrelevant.

The general idea involves overwriting a pointer that was previously returned by a call to malloc(), and that is subsequently passed to free(). This can lead to the linking of an arbitrary address into a fastbin. A further call to malloc() can result in this arbitrary address being used as a chunk of memory by the application. If the designer can control the applications use of the fake chunk, then it is possible to overwrite execution control data.

Assume that the designer has overflowed a pointer that is being passed to free(). The first problem that must be considered is exactly what the pointer should be overflowed with. Keep in mind that the ultimate goal of the House of Spirit is to allow the designer to overwrite some sort of execution control data by returning an arbitrary chunk to the application. Exactly what "execution control data" is doesn't particularly matter so long as overflowing it can result in execution being passed to a designer controlled memory location. The two most common examples that are suitable for use with the House of Spirit are function pointers and pending saved return addresses, which will herein be referred to as the "target".

In order to successfully apply the House of Spirit it is necessary to have a designer controlled word value at a lower address than the target. This word will correspond to the size field of the chunk header for the fakechunk passed to free(). This means that the overflowed pointer must be set to the address of the designer controlled word plus 4. Furthermore, the size of the fakechunk must be located no more than 64 bytes away from the target. This is because the default maximum data size for a fastbin entry is 64, and at least the last 4 bytes of data are required to overwrite the target.

There is one more requirement for the layout of the fakechunk data which will be described shortly. For the moment, assume that all of the above conditions have been met, and that a call to free() is made on the suitable fakechunk. A call to free() is handled by a wrapper function called public_fREe():

```
1
2 void
3 public_fREe(Void_t* mem)
4 {
5     mstate ar_ptr;
6     mchunkptr p;          /* chunk corresponding to mem */
```



```

7     ...
8     p = mem2chunk(mem);
9     if (chunk_IS_MMAPPED(p))
10    {
11        munmap_chunk(p);
12        return;
13    }
14    ...
15    ar_ptr = arena_for_chunk(p);
16    ...
17    _int_free(ar_ptr, mem);

```

In this situation `mem` is the value that was originally overflowed to point to a fakechunk. This is converted to the "corresponding chunk" of the fakechunk's data, and passed to `arena_for_chunk()` in order to find the corresponding arena. In order to avoid special treatment as an `mmap()` chunk, and also to get a sensible arena, the size field of the fakechunk header must have the `IS_MMAPPED` and `NON_MAIN_ARENA` bits cleared. To do this, the designer can simply ensure that the fake size is a multiple of 8. This would mean the internal function `_int_free()` is reached:

```

1 void _int_free(mstate av, Void_t* mem){
2     mchunkptr      p;           /* chunk corresponding to mem */
3     INTERNAL_SIZE_T size;       /* its size */
4     mfastbinptr*    fb;         /* associated fastbin */
5     ...
6     p = mem2chunk(mem);
7     size = chunksize(p);
8     ...
9     if ((unsigned long)(size) <= (unsigned long)(av->max_fast))
10    {
11        if (chunk_at_offset (p, size)->size <= 2 * SIZE_SZ
12            || __builtin_expect (chunksize (chunk_at_offset (p, size))
13                                >= av->system_mem, 0))
14        {
15            errstr = "free(): invalid next size (fast)";
16            goto errout;
17        }
18        ...
19        fb = &(av->fastbins[fastbin_index(size)]);
20        ...
21        p->fd = *fb;
22        *fb = p;
23    }
24 }

```

This is all of the code in `free()` that concerns the House of Spirit. The designer controlled value of `mem` is again converted to a chunk and the fake size value is extracted. Since `size` is designer controlled, the fastbin code can be triggered simply by ensuring that it is less than `av->max_fast`, which has a default of `64 + 8`. The final point of consideration in the layout of the fakechunk is the nextsize integrity tests.

Since the size of the fakechunk has to be large enough to encompass the target, the size of the nextchunk must be at an address higher than the target. The nextsize integrity tests must be handled for the fakechunk to be put in a fastbin, which means that there must be yet another designer controlled value at an address higher than the target.

The exact location of the designer controlled values directly depend on the size of the allocation request

that will subsequently be used by the designer to overwrite the target. That is, if an allocation request of N bytes is made (such that $N \leq 64$), then the designer's lower value must be within N bytes of the target and must be equal to $(N + 8)$. This is to ensure that the fakechunk is put in the right fastbin for the subsequent allocation request. Furthermore, the designer's upper value must be at $(N + 8)$ bytes above the lower value to ensure that the nextsize integrity tests are passed.

If such a memory layout can be achieved, then the address of this "structure" will be placed in a fastbin. The code for the subsequent `malloc()` request that uses this arbitrary fastbin entry is simple and need not be reproduced here. As far as `_int_malloc()` is concerned the fake chunk that it is preparing to return to the application is perfectly valid. Once this has occurred it is simply up to the designer to manipulate the application in to overwriting the target.

6 The House of Chaos

Virtuality is a dichotomy between the virtual adept and information, where the virtual adept signifies the infinite potential of information, and information is a finite manifestation of the infinite potential. The virtual adept is the conscious element of virtuality, the nature of which is to create and spread information. This is all that the virtual adept knows, and all that the virtual adept is concerned with.

When you talk to a very knowledgeable and particularly creative person, then you may well be talking to a hacker. However, you will never talk to a virtual adept. The virtual adept has no physical form, it exists purely in the virtual. The virtual adept may be contained within the material, contained within a person, but the adept itself is a distinct and entirely independent consciousness.

Concepts of ownership have no meaning to the virtual adept. All information belongs to virtuality, and virtuality alone. Because of this, the virtual adept has no concept of computer security. Information is invoked from virtuality by giving a request. In virtuality there is no level of privilege, no logical barrier between systems, no point of illegality. There is only information and those that can invoke it.

The virtual adept does not own the information it creates, and thus has no right or desire to profit from it. The virtual adept exists purely to manifest the infinite potential of information in to information itself, and to minimize the complexity of an information request in a way that will benefit all conscious entities. What is not information is not consequential to the virtual adept, not money, not fame, not power.

```
Am I a hacker? No.  
I am a student of virtuality.  
I am the witch malloc,  
I am the cult of the otherworld,  
and I am the entropy.  
I am Phantasmal Phantasmagoria,  
and I am a virtual adept.
```